

SDS 385: Homework 1

G. Paulon

September 12, 2016

Problem 1. Linear regression

(A) Let us write the WLS objective (or cost) function

$$l(\beta) = \sum_{i=1}^N \frac{w_i}{2} (y_i - x_i^T \beta)^2$$

in a matricial form. We get

$$\begin{aligned} l(\beta) &= \frac{1}{2} (y - X\beta)^T W (y - X\beta) \\ &= \frac{1}{2} (y^T - (X\beta)^T) W (y - X\beta) \\ &= \frac{1}{2} (y^T W y - y^T W X \beta - (X\beta)^T W y + (X\beta)^T W X \beta) \end{aligned}$$

Let us remark here that the second and the third term of the sum can be collected in one single term. In fact

$$(y^T W X \beta)^T = (X \beta)^T W y,$$

i.e. the first is the transposed of the second but, since they are real numbers, they are also equal. Hence we have

$$l(\beta) = \frac{1}{2} (y^T W y - 2\beta^T X^T W y + \beta^T X^T W X \beta).$$

Since we are trying to minimize a convex function, we can just find its stationary points. To do this, we first need to calculate the gradient of the WLS objective function equalizing it to 0. Thus,

$$\begin{aligned} \nabla l(\hat{\beta}) &= 0 \\ \Leftrightarrow \frac{1}{2} (-2X^T W y + 2X^T W X \hat{\beta}) &= 0 \\ \Leftrightarrow (X^T W X) \hat{\beta} &= X^T W y, \end{aligned}$$

which is the set of desired equations (also known as normal equations).

(B) Solving the linear system requires the matrix $X^T W X$ to be nonsingular (and therefore invertible). The matrix $X^T W X$ is positive definite, and therefore nonsingular, in case X has full rank (i.e. the columns of X are linearly independent). The normal equations can be easily solved by inverting the matrices, which leads to

$$\hat{\beta} = (X^T W X)^{-1} X^T W y.$$

The classical inversion method, however, is not recommended as its computational complexity grows rapidly as the dimension of the matrices increases. Moreover, the inversion of a matrix is not stable numerically speaking, as small perturbations on the original matrix

can produce high differences in the inverse matrices (i.e. the numerical solution does not approach the exact one).

For these reasons, we present here two other approaches that are used to deal with the solution of linear systems: the QR and the Cholesky factorizations.

- The QR factorization consists in finding two matrices Q (orthogonal) and R (upper triangular) such that a matrix A can be written as their product. In our case, we want to write the QR factorization of $W^{\frac{1}{2}}X$, i.e. $W^{\frac{1}{2}}X = QR$. Therefore we get

$$\begin{aligned}(X^T W X) \hat{\beta} &= X^T W y \\ \Leftrightarrow (W^{\frac{1}{2}} X)^T (W^{\frac{1}{2}} X) \hat{\beta} &= (W^{\frac{1}{2}} X)^T W^{\frac{1}{2}} y \\ \Leftrightarrow (QR)^T QR \hat{\beta} &= (QR)^T W^{\frac{1}{2}} y \\ \Leftrightarrow R^T Q^T QR \hat{\beta} &= R^T Q^T W^{\frac{1}{2}} y \\ \Leftrightarrow R \hat{\beta} &= Q^T W^{\frac{1}{2}} y.\end{aligned}$$

This system can be easily solved because R is upper triangular.

In our case, the pseudocode can be summarized as follows.

Algorithm 1 QR solver

- 1: **procedure** MY_QR(X, y, W)
 - 2: find the factors Q, R , s.t. $W^{\frac{1}{2}}X = QR$
 - 3: compute the right-term: $b \leftarrow Q^T W^{\frac{1}{2}} y$
 - 4: solve the linear system: $\hat{\beta} \leftarrow R^{-1} b$
-

- The Cholesky factorization can be found for symmetric and positive definite matrices, and it consists in factorizing the original matrix as $A = RR^T$, where R is a lower triangular matrix.

In our case, the pseudocode can be summarized as follows.

Algorithm 2 Cholesky solver

- 1: **procedure** MY_CHOL(X, y, W)
 - 2: find the factor R , s.t. $X^T W X = RR^T$
 - 3: solve the first linear system: $q \leftarrow R^{-1} X^T W y$
 - 4: solve the second linear system: $\hat{\beta} \leftarrow R^{-T} q$
-

(C) We here compare the three methods used to solve the linear system of the normal equations. Let us point out that in each of the three algorithms we exploited the fact that W is a diagonal matrix, optimizing the matricial product $X^T W$ by simply multiplying the first element of the diagonal of W by the first column of X^T , and so on. Moreover, since the matrix $X^T W X$ is

symmetric, one can exploit this information to calculate it by using R's `crossprod(W1/2X)` function.

The performances of the three algorithms are displayed in Figure 1. Four test cases were analysed: the values of N and P range from a minimum of $(N, P) = (100, 10)$ to a maximum of $(N, P) = (3000, 1000)$. As one can see, Cholesky method is the fastest one. The inversion is slower than Cholesky method and it is also more unstable than every other method, and therefore it has to be avoided. The QR factorization, despite being the slowest method, is preferable when numerical issues can occur. In fact, this method is the most robust one.

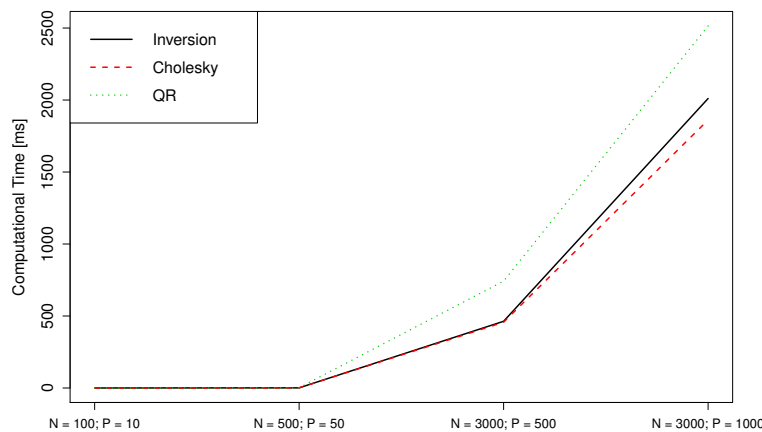


Figure 1: Comparison between the performances of the three algorithms for different dimensions of the data and of their features.

- (D) In the following we use Cholesky method instead of the QR factorization, as the first is faster and now the primary interest is the performance of the algorithms. If X is a highly sparse matrix, one can exploit appropriate routines in R language in order to make the computation faster.

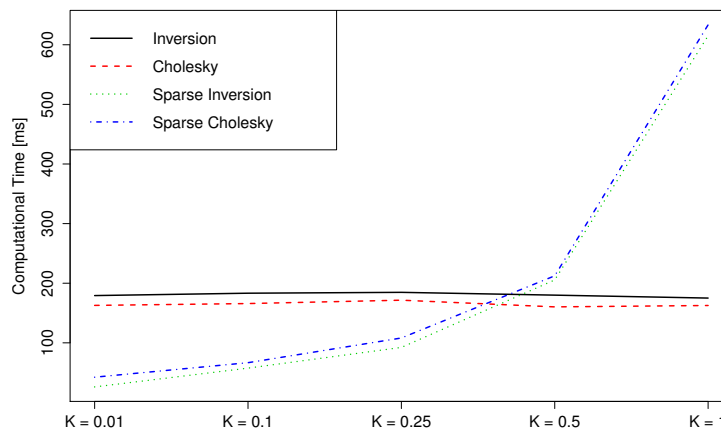


Figure 2: Comparison between the performances of the three algorithms for different sparsity levels K .

Let us take, as an example, $(N, P) = (1000, 500)$ and let us compare the results for different levels of sparsity K of the matrix X . As one can see in Figure 2, when the features matrix is dense, both the Cholesky and the inversion methods outperform the sparse methods, as these latter involve a first stage in which X is converted in a sparse format. However, as K increases, the sparse methods become significantly faster than the others. This is due to the fact that matricial products are more efficient in a sparse format and a lot of useless computations is avoided.

Problem 2. Generalized linear models

(A) The negative log-likelihood for the problem of logistic regression can be written as

$$\begin{aligned} l(\beta) &= -\log \left\{ \prod_{i=1}^N p(y_i|\beta) \right\} \\ &= -\sum_{i=1}^N \left\{ \log \binom{m_i}{y_i} + y_i \log(w_i) + (m_i - y_i) \log(1 - w_i) \right\}. \end{aligned} \quad (1)$$

Remark that for extreme values of $x_i^T \beta$ the probabilities w_i approach 0 and 1, making the expression subject to numerical issues. This problem has been solved with a naive approach, by summing a small constant in the argument of the logarithm. More sophisticated approaches are possible, such as giving an asymptotical approximation for (1) that can handle these extreme cases. As an alternative, one could use a Lasso or Ridge regression, introducing a penalty term for β .

The gradient of (1) is

$$\begin{aligned} \nabla l(\beta) &= -\sum_{i=1}^N \left\{ y_i \frac{1}{w_i} \frac{x_i e^{-x_i^T \beta}}{(1 + e^{-x_i^T \beta})^2} - (m_i - y_i) \frac{1}{1 - w_i} \frac{x_i e^{-x_i^T \beta}}{(1 + e^{-x_i^T \beta})^2} \right\} \\ &= -\sum_{i=1}^N \left\{ y_i \frac{x_i e^{-x_i^T \beta}}{1 + e^{-x_i^T \beta}} - m_i \frac{x_i}{1 + e^{-x_i^T \beta}} + y_i \frac{x_i}{1 + e^{-x_i^T \beta}} \right\} \\ &= -\sum_{i=1}^N \left\{ y_i x_i - m_i \frac{x_i}{1 + e^{-x_i^T \beta}} \right\} \\ &= -\sum_{i=1}^N \{x_i (y_i - m_i w_i)\} \\ &= -X^T (y - mw). \end{aligned}$$

(B) The gradient descent method has been implemented for a fixed step size $\alpha = 0.01$, which is a sufficiently small value to ensure convergence. As a stopping criterion, we used the relative decrement of the negative log-likelihood. In other terms, we defined the quantity

$$\delta = \frac{l(\beta^{(k)}) - l(\beta^{(k+1)})}{|l(\beta^{(k)}) + \varepsilon|}$$

and the algorithm stops when δ is smaller than a certain threshold (in our case set to 10^{-10}). Remark that this criterion is the definition of convergence from a numerical point of view, i.e. when the objective function does not change anymore. Other choices are valid: one could, for example, use the absolute value of the relative difference of the parameters β over iterations. A different approach would consist in exploiting an analytical criterion, such as stopping the algorithm when the L_2 -norm of the gradient is below a certain threshold.

Gradient descent is known to suffer to the problem of local optima. However, in this case we are trying to minimize a convex function, which ensures that the algorithm converges to the unique global minimum. In fact, the Hessian matrix is positive definite, as we show in the next part.

We tested the validity of this method on a real dataset containing the classification of a breast cell. We dispose of 10 features for X , and of $n = 569$ data points. In Figure 3 the values of the negative log-likelihood are displayed. The gradient descent converges, in this case, after around 60000 iterations.




Figure 3: Negative log-likelihood of the data as the gradient descent converges; the x-axis is reported on the log-scale.

- (C) In order to write the second-order Taylor approximation, we first need to compute the Hessian of the negative log-likelihood, that is

$$\begin{aligned}\nabla^2 l(\beta) &= \nabla \left(- \sum_{i=1}^N x_i (y_i - m_i w_i) \right) \\ &= \left(\sum_{i=1}^N \nabla(m_i x_{i1} w_i); \dots; \sum_{i=1}^N \nabla(m_i x_{ip} w_i) \right).\end{aligned}$$

By exploiting the fact that $\frac{\partial w_i}{\partial \beta_j} = x_{ij}w_i(1 - w_i)$, we obtain,

$$\begin{aligned}\nabla^2 l(\beta) &= \begin{pmatrix} \sum_{i=1}^N m_i x_{i1} x_{i1} w_i(1 - w_i) & \dots & \sum_{i=1}^N m_i x_{ip} x_{i1} w_i(1 - w_i) \\ \vdots & \ddots & \vdots \\ \sum_{i=1}^N m_i x_{i1} x_{ip} w_i(1 - w_i) & \dots & \sum_{i=1}^N m_i x_{ip} x_{ip} w_i(1 - w_i) \end{pmatrix} \\ &= X^T W X,\end{aligned}$$

where the matrix W is defined as $W = \text{diag}\{m_1 w_1(1 - w_1); \dots; m_N w_N(1 - w_N)\}$. Let us remark that, if W has positive entries (which is the case), the matrix $X^T W X$ is positive definite.

The second order Taylor expansion is then

$$\begin{aligned}\tilde{l}(\beta) &= l(\beta_0) + \nabla l(\beta_0)^T (\beta - \beta_0) + \frac{1}{2} (\beta - \beta_0)^T \nabla^2 l(\beta_0) (\beta - \beta_0) \\ &= l(\beta_0) - (y - mw)^T X (\beta - \beta_0) + \frac{1}{2} (\beta - \beta_0)^T X^T W X (\beta - \beta_0) \\ &= l(\beta_0) - (y - mw)^T X \beta + (y - mw)^T X \beta_0 \\ &\quad + \frac{1}{2} \beta^T X^T W X \beta + \frac{1}{2} \beta_0^T X^T W X \beta_0 - \beta_0^T X^T W X \beta \\ &= l(\beta_0) + (y - mw)^T X \beta_0 + \frac{1}{2} \beta_0^T X^T W X \beta_0 \\ &\quad - (y - mw + W X \beta_0)^T X \beta + \frac{1}{2} (X \beta)^T W (X \beta) \\ &= \frac{1}{2} (z - X \beta)^T W (z - X \beta) + c,\end{aligned}$$

where

$$\begin{aligned}z &= -W^{-1}(y - mw + W X \beta_0) = W^{-1}(mw - y) - X \beta_0; \\ c &= l(\beta_0) + (y - mw)^T X \beta_0 + \frac{1}{2} \beta_0^T X^T W X \beta_0 \\ &\quad - \frac{1}{2} (y - mw + W X \beta_0)^T W^{-1} (y - mw + W X \beta_0).\end{aligned}$$

(D) Newton's method has successfully been implemented. As a stopping criterion, we used the same as before.

We tested the validity of this method on the same real dataset as before. In Figure 4 the values of the negative log-likelihood are displayed. The gradient descent converges, in this case, after 11 iterations. In Listing 1 we can see that the estimates using R's built-in `glm` function and the two methods we implemented are the same.

(E) As we have seen, Newton's method is significantly faster than gradient descent in terms of the total number of iterations that are needed to reach convergence. However, be aware that each iteration of Newton's method is computationally more intense, as it involves the




Figure 4: Negative log-likelihood of the data as Newton's method converges.

1		Beta0	Beta1	Beta2	Beta3	Beta4	Beta5	Beta6	Beta7
2	Glm	0.4870168	-7.221851	1.654756	-1.737630	14.00485	1.074953	-0.07723455	0.6751231
3	Gradient	0.4747997	-6.597393	1.655226	-2.311252	13.90645	1.069264	-0.04381625	0.6801968
4	Newton	0.4870168	-7.221851	1.654756	-1.737630	14.00485	1.074953	-0.07723455	0.6751231
5									
6		Beta8	Beta9	Beta10					
7	Glm	2.592874	0.4462563	-0.4824842					
8	Gradient	2.600741	0.4454228	-0.4847644					
9	Newton	2.592874	0.4462563	-0.4824842					

Listing 1: Comparison of the estimates using three different algorithms

computation of the Hessian matrix and, moreover, the resolution of a linear system. Overall, Newton's method converges more rapidly as it is able to fully exploit the information provided by the second derivatives of the function we are trying to minimize, that is, the information concerning the curvature of the objective function. In fact, choosing the descent direction to be simply the vector orthogonal to the level curves can often lead to a slower convergence. The inverse of the Hessian matrix has the role to correct the gradient direction.

Newton's method is reliable when the difference between the true function and its quadratic approximation is not too large. However, cases when the objective function is not twice differentiable may arise. In these cases, the Hessian is not defined. Moreover, dealing with big datasets (let us say, for example, with 1000 predictors) implies that the Hessian can be very large and its computation can be cumbersome.

In conclusion, Newton's method does not scale as well as gradient descent on big data.

Appendix A

R code

```
1 # Functions used to solve the linear system
2 #  $\text{Beta.hat} = (X^T W X)^{-1} X^T W Y$ 
3 # whose solution minimizes
4 #  $\frac{1}{2}(Y - X \text{Beta})^T W (Y - X \text{Beta})$ 
5
6 # Optimal function calculating WLS via inverting the matrix and optimizing the product with the
  diagonal matrix
7 my.inv <- function(X, y, W){
8   wx <- diag(W)^(1/2) * X
9   wy <- diag(W)^(1/2) * y
10  beta.hat <- solve(crossprod(wx), crossprod(wx, wy))
11
12  return(beta.hat)
13 }
14
15 # Optimal function calculating WLS via Cholesky decomposition and optimizing the product with the
  diagonal matrix
16 my.chol <- function(X, y, W){
17   wx <- diag(W)^(1/2) * X
18   wy <- diag(W)^(1/2) * y
19   R = chol(crossprod(wx))
20   u = forwardsolve(t(R), crossprod(wx, wy))
21   beta.hat = backsolve(R, u)
22
23   return(beta.hat)
24 }
25
26 # Optimal function calculating WLS via QR factorization
27 my.QR <- function(X, y, W){
28   P <- ncol(X)
29   wx <- diag(W)^(1/2) * X
30   wy <- diag(W)^(1/2) * y
31   QR <- qr(wx)
32
33   qty = qr.qty(QR, wy)
34   beta.hat = backsolve(QR$qr, qty)
35
36   return(beta.hat)
37 }
38
```

```

39 # Optimal function calculating WLS via inverting the matrix and exploiting the sparsity of the
    matrix X
40 my.invsparse <- function(X, y, W){
41   X = Matrix(X, sparse=TRUE)
42   wx <- diag(W)^(1/2) * X
43   wy <- diag(W)^(1/2) * y
44   beta.hat <- solve(crossprod(wx), crossprod(wx, wy))
45
46   return(beta.hat)
47 }
48
49 # Optimal function calculating WLS via Cholesky decomposition and exploiting the sparsity of the
    matrix X
50 my.cholsparse <- function(X, y, W){
51   X = Matrix(X, sparse=TRUE)
52   wx <- diag(W)^(1/2) * X
53   wy <- diag(W)^(1/2) * y
54   R = chol(crossprod(wx))
55   u = forwardsolve(t(R), crossprod(wx, wy))
56   beta.hat = backsolve(R, u)
57
58   return(beta.hat)
59 }

```

Listing A.1: Functions used to solve the linear system for the weighted least squares problem

```

1 # Compute the probabilities associated with the logit model
2 comp.wi <- function(X, beta){
3   wi <- 1/(1+exp(-X %*% beta))
4   return(wi)
5 }
6
7 # Compute the negative log-likelihood of the logit model
8 log.lik <- function(beta, y, X, mi){
9   ll <- array(NA, dim = N)
10  wi <- comp.wi(X, beta)
11  ll <- -sum(y * log(wi + 1E-4) + (mi - y) * log(1 - wi + 1E-4))
12  return(ll)
13 }
14
15 # Compute the gradient of the negative log-likelihood of the logit model
16 grad.loglik <- function(beta, y, X, mi){
17   grad.ll <- array(NA, dim = length(beta))
18   wi <- comp.wi(X, beta)
19   grad.ll <- -apply(X * as.numeric(y - mi * wi), 2, sum)
20   return(grad.ll)
21 }
22
23 # Compute the hessian of the negative log-likelihood of the logit model
24 hessian.loglik <- function(beta, y, X, mi){
25   wi <- comp.wi(X, beta)
26   W.12 <- sqrt(mi * wi * (1 - wi))
27   hes.ll <- crossprod(as.numeric(W.12) * X)
28   return(hes.ll)
29 }
30
31 # Function for the gradient descent of the logit model

```

```

32 gradient.descent <- function(y, X, mi, beta0, maxiter, alpha, tol){
33
34   betas <- array(NA, dim=c(maxiter, length(beta0)))
35   betas[1,] <- beta0 # Initial guess
36   ll <- array(NA, dim = maxiter)
37   ll[1] <- log.lik(betas[1,], y, X, mi)
38
39   for (iter in 2:maxiter){
40     gradient <- grad.loglik(betas[iter-1,], y, X, mi)
41     betas[iter,] <- betas[iter-1,] - alpha * gradient
42     ll[iter] <- log.lik(betas[iter,], y, X, mi)
43
44     # Convergence check
45     if ((ll[iter-1] - ll[iter]) / (ll[iter-1] + 1E-3) < tol){
46       cat('Algorithm has converged after', iter, 'iterations')
47       ll <- ll[1:iter]
48       betas <- betas[1:iter,]
49       break;
50     }
51
52     else if (iter == maxiter & (ll[iter-1] - ll[iter]) / (ll[iter-1] + 1E-2) >= tol){
53       print('WARNING: algorithm has not converged')
54       break;
55     }
56   }
57   return(list("ll" = ll, "beta" = betas[iter,]))
58 }
59
60 # Function for the Newton method of the logit model
61 newton.descent <- function(y, X, mi, beta0, maxiter, tol){
62
63   betas <- array(NA, dim=c(maxiter, length(beta0)))
64   betas[1,] <- beta0 # Initial guess
65   ll <- array(NA, dim = maxiter)
66   ll[1] <- log.lik(betas[1,], y, X, mi)
67
68   for (iter in 2:maxiter){
69     hessian <- hessian.loglik(betas[iter-1,], y, X, mi)
70     gradient <- grad.loglik(betas[iter-1,], y, X, mi)
71
72     # Solve the linear system  $H^{(-1)} \cdot \text{Grad}$  in order to find the direction
73     direct <- QR.solver(hessian, gradient)
74     betas[iter,] <- betas[iter-1,] - direct
75     ll[iter] <- log.lik(betas[iter,], y, X, mi)
76
77     # Convergence check
78     if ((ll[iter-1] - ll[iter]) / (ll[iter-1] + 1E-3) < tol){
79       cat('Algorithm has converged after', iter, 'iterations')
80       ll <- ll[1:iter]
81       betas <- betas[1:iter,]
82       break;
83     }
84
85     else if (iter == maxiter & (ll[iter-1] - ll[iter]) / (ll[iter-1] + 1E-2) >= tol){
86       print('WARNING: algorithm has not converged')
87       break;
88     }

```

```
89  }  
90  return(list("ll" = ll, "beta" = betas[iter,]))  
91 }
```

Listing A.2: Functions implementing descent methods for the logistic regression problem