

# SDS 385: Homework 4

G. Paulon

October 7, 2016



### Problem 1. Improving SGD for logistic regression

- (A) We first implement the stochastic gradient descent coupled with backtracking for logistic regression. This algorithm uses minibatches of the data in order to pick the correct step size for a certain number of iterations. The step size is refreshed every  $k$  iterations, which leads to a natural decay over time. The code is reported in Listing A.1. The results of this algorithm applied to the cancer dataset are not as promising as the results obtained with the adaptive gradient descent. For this reason, in the following we will discuss an efficient implementation and a big-data application only of the second algorithm.
- (B) The Adaptive Gradient Descent has been implemented. Let us describe the updating rules for this algorithm, which uses a diagonal approximation of the Hessian at every iteration. Let us first recall that the negative log-likelihood of the  $i^{th}$  data point is given by

$$\begin{aligned}
 l_i(\alpha, \beta) &\propto -y_i \log(w_i) - (m_i - y_i) \log(1 - w_i) - \sum_{j=1}^p |\beta_j| \\
 &= -m_i \log(1 - w_i) - y_i \log\left(\frac{w_i}{1 - w_i}\right) - \sum_{j=1}^p |\beta_j| \\
 &= m_i \log(1 + \exp(\alpha + \mathbf{x}_i^T \beta)) - y_i(\alpha + \mathbf{x}_i^T \beta) - \sum_{j=1}^p |\beta_j|
 \end{aligned}$$

where we added a penalization term that forces sparsity in the problem. Let us remark that we split the intercept and the other covariates. This is done for notation purposes, as the intercept is not penalized in the last term of the sum.

The gradient with respect to  $\alpha$  of this contribution is

$$\begin{aligned}
 \nabla_{\alpha} l_i(\alpha, \beta) &\propto m_i \frac{e^{\alpha + \mathbf{x}_i^T \beta}}{1 + e^{\alpha + \mathbf{x}_i^T \beta}} - y_i \\
 &= \hat{y}_i - y_i,
 \end{aligned}$$

where

$$\hat{y}_i = m_i \frac{e^{\alpha + \mathbf{x}_i^T \beta}}{1 + e^{\alpha + \mathbf{x}_i^T \beta}}.$$

The gradient with respect to  $\beta$  of the individual negative log-likelihood is

$$\begin{aligned}
 \nabla_{\beta_j} l_i(\alpha, \beta) &\propto m_i \frac{x_{ij} e^{\alpha + \mathbf{x}_i^T \beta}}{1 + e^{\alpha + \mathbf{x}_i^T \beta}} - x_{ij} y_i - \lambda \frac{\beta_j}{|\beta_j|} \\
 &= x_{ij}(\hat{y}_i - y_i) - \lambda \frac{\beta_j}{|\beta_j|}.
 \end{aligned}$$

Adagrad still has a base learning rate  $\eta$ , but this is multiplied with the elements of a vector  $\mathbf{g}$ , which is given by

$$\mathbf{g}_j^{(T)} = \sum_{t=1}^T \left( \nabla_{\beta_j} l_i(\alpha, \beta^{(t)}) \right)^2.$$

The vector  $\mathbf{g}$  is monotonically increasing and it can be interpreted as the historical gradient, since it adds at every iteration the square of the gradient contribution. The update step is

$$\beta_j^{(t+1)} = \beta_j^{(t)} - \frac{\eta}{\sqrt{g_j^{(T)}}} \nabla_{\beta_j} l_i(\alpha, \beta^{(t)})$$

and as one can see, the step size decays naturally over time.

The code is illustrated in Listing A.2.

The algorithm has been applied to the usual dataset containing a cancer classification problem. The speed-up is remarkable if compared to the same version of the algorithm implemented in R language. The convergence results are displayed in Figure 1.

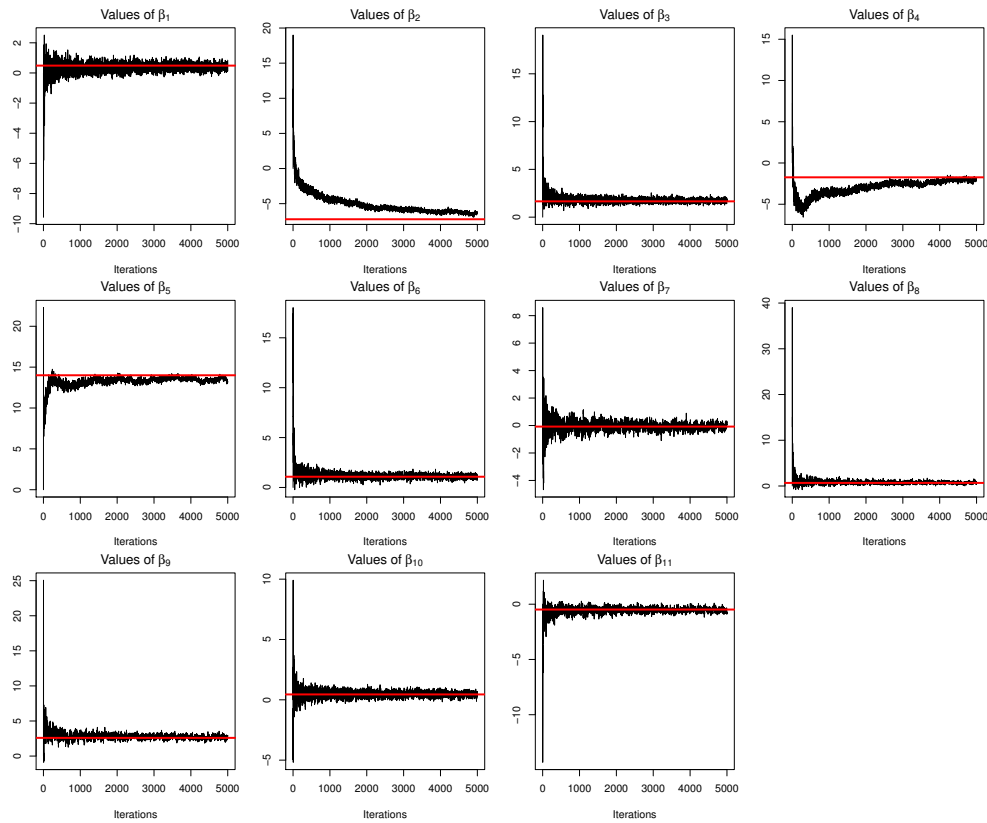


Figure 1: Results of the Adagrad method. The red lines denote the results obtained with Newton method.

## Problem 2. Putting it all together on some biggish data

The algorithm has been improved by adding support for sparse matrices and it has been applied to a large (and sparse) dataset whose goal is to detect malicious URLs in web traffic.

The optimal value of the regularizer parameter  $\lambda$  has been chosen by 5-fold cross validation. The estimate of the test error (prediction accuracy) in this case is around 98.0%, which is a remarkable result.

	$\lambda = 0$	$\lambda = 10^{-10}$	$\lambda = 10^{-9}$	$\lambda = 10^{-8}$	$\lambda = 10^{-7}$	$\lambda = 10^{-6}$	$\lambda = 10^{-5}$	$\lambda = 10^{-4}$
Mean	0.9856	0.9856	0.9855	0.9853	0.9841	0.9787	0.9724	0.9626
Std. dev.	0.0006	0.0005	0.0005	0.0006	0.0008	0.0014	0.0019	0.0028

Table 1: Cross-validation test error estimates for several values of the regularization parameter  $\lambda$ .

The details of the algorithm are explained in Listing A.2. However, we discuss the main features of this implementation. In order to make the algorithm as efficient as possible, we added:

- support for **sparse matrices**  $X$ : in fact, computing the dot product  $x_i^T \beta$  at every iteration can be efficiently done only by ignoring all the zero elements of  $x_i^T$ ;
- **column-major format** for the matrix  $X$ . Since C++ stores matrices in a column-major format (that is, contiguous cells are stored in memory following the vertical direction), we extract at each iteration the vector of the  $i^{th}$  observation by the column  $i$  of  $X$ . To do so, the original  $X$  matrix has been transposed before the algorithm starts;
- use of the **iterators**. Iterators are efficient pointers that allow to loop very quickly over matrices and vector (most of all, if they are sparse) by pointing to the addresses of the non-empty cells;
- **lazy update** of the penalization when the parameter  $\beta$  has not been updated for some iterations. In particular, in this version of the implementation we chose to use a  $L^1$  penalization because it is more likely to fit well our problem, which has around 2000000 features, by shrinking some features towards 0. Let us suppose that for  $k$  consecutive iterations (that is, for  $k$  consecutive observations) the feature  $\beta_j$  is not updated because  $x_{ij} = 0$ . The gradient contribution of the negative log-likelihood is 0 and so  $\beta_j$  is not updated. However, we should still update the parameter because of the penalty term, which is never 0. In order to avoid to do this computation for every iteration (which would compromise the computational efficiency of having a sparse  $X$ ) we update the cumulative penalty for  $\beta_j$  only when is updated.

If  $k$  updates have been skipped from step  $t+1$  to step  $t+k$  (that is,  $\beta_j^{(t)}$  is the last one updated and  $\beta_j^{(t+k+1)}$  is the first one to be updated again), then:

$$\begin{aligned}
\beta_j^{(t+1)} &= \beta_j^{(t)} - \frac{\eta}{\sqrt{g_j^{(T)}}} \lambda \text{sign}(\beta_j^{(t)}) \\
\beta_j^{(t+2)} &= \beta_j^{(t+1)} - \frac{\eta}{\sqrt{g_j^{(T)}}} \lambda \text{sign}(\beta_j^{(t+1)}) = \beta_j^{(t)} - \frac{\eta}{\sqrt{g_j^{(T)}}} 2\lambda \text{sign}(\beta_j^{(t)}) \\
&\vdots \\
\beta_j^{(t+k)} &= \beta_j^{(t)} - \frac{\eta}{\sqrt{g_j^{(T)}}} k\lambda \text{sign}(\beta_j^{(t)})
\end{aligned}$$

where we assumed that  $\beta_j$  does not change its sign during the  $k$  iterations, and that the historical gradient  $g_j^{(T)}$  is constant as well. Afterwards,  $\beta_j^{(t+k+1)}$  is updated with the usual rule. In order to make our assumptions more reasonable,

$$\beta_j^{(t+k)} = \text{sign}(\beta_j^{(t)}) \left( 0; |\beta_j^{(t)}| - \frac{\eta k \lambda}{\sqrt{g_j^{(T)}}} \right)_+,$$

which consists in setting  $\beta_k^{(t+k)}$  to 0 if it would change sign during the lazy update.

# Appendix A

## R and C++ code

```
1 SGD.linesearch <- function(y, X, mi, beta0, maxiter = 100000, tol = 1E-8){
2   # Function for the gradient descent of the logit model coupled with backtracking
3   # -----
4   # Args:
5   #   - y: response vector (length n)
6   #   - X: matrix of the features (n*p)
7   #   - mi: vector of the number of trials, always 1 in the logit framework (length n)
8   #   - beta0: initial regression parameters (length p)
9   #   - maxiter: number of maximum iterations the algorithm will perform if not converging
10  #   - tol: tolerance threshold to convergence
11  # Returns:
12  #   - ll: values of the log-likelihood for every iteration
13  #   - beta: final optimal regression parameters
14  #   - alpha: selected step size for every iteration
15  # -----
16  N <- dim(X)[1]
17  P <- length(beta0)
18
19  # Select the updating time and the size of the minibatches
20  train.epoch <- floor(maxiter/1000)
21  size.minibatch <- floor(N/10)
22
23  # Sample the data points to calculate the gradient
24  idx <- sample(1:N, maxiter, replace = TRUE)
25  # Sample the indexes for the minibatch
26  minibatches <- matrix(sample(1:N, 1000*size.minibatch, replace = TRUE), 1000, size.minibatch)
27  idx.minibatch <- minibatches[1, ]
28
29  betas <- array(NA, dim=c(maxiter, length(beta0)))
30  betas[1, ] <- beta0 # Initial guess
31
32  av.ll <- array(NA, dim = maxiter)
33  ll <- log.lik(betas[1, ], y[idx[1]], matrix(X[idx[1],], nrow=1), mi[idx[1]])
34  av.ll[1] <- ll
35
36  for (iter in 2:maxiter){
37
38    # Choose the direction according to the gradient of the single individual idx[i]
39    grad <- grad.loglik(betas[iter-1,], y[idx[iter]], matrix(X[idx[iter], ],nrow=1), mi[idx[iter]
    ])
```

```

40  dir <- - grad
41  # if we need to update the optimal step
42  if ((iter == 2) | (iter %% train.epoch == 0)){
43    idx.minibatch <- minibatches[ceiling(iter/train.epoch), ]
44    grad.minibatch <- grad.loglik(betas[iter-1, ], y[idx.minibatch], X[idx.minibatch, ], mi[idx
      .minibatch]) / length(idx.minibatch)
45    alpha <- optimal.step(betas[iter-1, ], -grad.minibatch, grad.minibatch, y[idx[iter]],
      matrix(X[idx[iter], ], nrow=1), mi[idx[iter]])
46  }
47
48  # Update Beta parameters
49  betas[iter, ] <- betas[iter-1, ] + alpha * dir
50  # Compute log-likelihood and average log-likelihood
51  ll <- log.lik(betas[iter, ], y[idx[iter]], matrix(X[idx[iter], ], nrow=1), mi[idx[iter]])
52  av.ll[iter] <- ((av.ll[iter-1])*(iter-1) + ll)/iter
53
54  # Convergence check
55  if (abs(av.ll[iter-1] - av.ll[iter]) / (av.ll[iter-1] + 1E-10) < tol){
56    av.ll <- av.ll[1:iter]
57    betas <- betas[1:iter, ]
58    break;
59  }
60
61  else if (iter == maxiter & abs(av.ll[iter-1] - av.ll[iter])/(av.ll[iter-1] + 1E-10) >= tol){
62    break;
63  }
64  }
65  return(list("ll" = av.ll, "beta" = betas))
66  }

```

*Listing A.1: Function implementing the SGD coupled with backtracking.*

```

1  #define ARMA_64BIT_WORD
2  // [[Rcpp::depends(RcppEigen)]]
3  #include <RcppEigen.h>
4  #include <iostream>
5  #include <algorithm>
6  #include <string>
7
8  using namespace Rcpp;
9
10 //using Eigen::Map;
11 //using Eigen::MatrixXd;
12 //using Eigen::MatrixXi;
13 using Eigen::VectorXd;
14 //using Eigen::VectorXi;
15 using Eigen::SparseVector;
16
17 typedef Eigen::MappedSparseMatrix<double> MapMatd;
18 // typedef Map<MatrixXi> MapMati;
19 // typedef Map<VectorXd> MapVecd;
20 // typedef Map<VectorXi> MapVeci;
21
22
23 // Function to compute the sign of a generic type
24 template <typename T> int sgn(T val) {
25   return (T(0) < val) - (val < T(0));

```



```

26 }
27
28
29 // [[Rcpp::export]]
30 SEXP Ada_Grad(MapMatd& X, VectorXd& y, VectorXd& m, VectorXd& beta0, double eta = 1.0, unsigned
    int npass = 1, double lambda = 0.0, double weight = 0.01){
31 // - X:      design matrix stored in sparse column-major format. That is, we transposed the
    matrix outside the function (in R).
32 //      The features for the observation i is stored in column i
33 // - y:      response vector
34 // - m:      vector of sample sizes
35 // - beta0:   initial guess values for beta
36 // - eta:     master step-size
37 // - npass:   number of times we go over the dataset
38 // - lambda:  penalization of the Lasso regression (L1 penalty)
39 // - weight:  weight for the computation of the negative log-likelihood (high weight gives
    importance at every single contribution, small weight smoothes the
40 //      negative log-likelihood function).
41
42 unsigned int N = X.cols();
43 unsigned int P = X.rows();
44
45 // x is the generic column of X that will be extracted at each iteration
46 SparseVector<double> x(P);
47
48 // Initialize parameters
49 double psi, epsi, yhat, delta, h;
50 double this_grad = 0.0;
51 double newbeta, penalty;
52 unsigned int j,k;
53
54 // Initialize parameters alpha and beta
55 double alpha = 0.0;
56 VectorXd beta(P);
57
58 // Initialize historical gradients (cumulative)
59 double hist_int = 0.0; // historical gradient of the intercept term
60 VectorXd hist_grad(P);
61 for (int j = 0; j < P; j++){
62     hist_grad(j) = 1e-3;
63     beta(j) = beta0(j);
64 }
65
66 // How long has it been since the last update of each feature?
67 NumericVector last_update(P, 0.0);
68
69 // negative log likelihood for assessing fit
70 double loglik_avg = 0.0;
71 NumericVector loglik_vec(npass*N, 0.0);
72
73
74 // Outer loop: number of passes over data set
75 k = 0; // global iteration counter
76 for(unsigned int pass = 0; pass < npass; pass++) {
77
78     // Loop over each observation (columns of X)
79     for(unsigned int i = 0; i < N; i++) {

```

```

80
81 // Compute linear predictor and yhat
82 x = X.innerVector(i); // efficient way to extract a column from the matrix X
83 psi = alpha + x.dot(beta); // the function .dot() automatically ignores the zero
    elements of x
84 epsi = exp(psi);
85 yhat = m(i) * epsi/(1.0 + epsi); // MLE of y
86
87 // (1) Update log-likelihood weighted average
88 loglik_avg = (1.0 - weight) * loglik_avg + weight * (m(i) * log(1 + epsi) - y(i) * psi);
89 loglik_vec(k) = loglik_avg;
90
91 // (2) Update intercept
92 delta = y(i) - yhat; // GRADIENT with respect to the intercept (the
    intercept component)
93 hist_int += delta * delta; // update historical gradient of the intercept
94 alpha += (eta/sqrt(hist_int)) * delta; // compute the update of the intercept
95
96 // (3) Update beta: iterate over the ACTIVE features for this instance
97 for (SparseVector<double>::InnerIterator it(x); it; ++it) {
98
99     // the .index() function extracts the x index of iterator currently pointing at x
100     j = it.index();
101
102     // STEP (a): aggregate all the penalty-only updates since the last time we updated this
        feature.
103     // This is a form of lazy updating in which we approximate all the "penalty-only" updates
        at once.
104     double skip = k - last_update(j); // how many penalization updates did we skip for
        this feature?
105     h = sqrt(hist_grad(j));
106     penalty = skip * eta/h; // penalize the skip previous updates
107
108     // We can find the corresponding beta applying the penalty. Let us remark that if beta has
        changed sign, we set it to 0. This is reasonable
109     beta(j) = sgn(beta(j)) * fmax(0.0, fabs(beta(j)) - penalty*lambda);
110     // For the L2 penalty, we can change this line to: 2*lambda*skip*beta(j), which is an
        approximation because beta is changing at each iteration
111
112     // Update the last-update vector
113     last_update(j) = k;
114
115     // STEP (b): Now we compute the regular update for this observation.
116     this_grad = delta*it.value();
117
118     // update the historical gradient for this feature
119     hist_grad(j) += this_grad*this_grad;
120
121     h = sqrt(hist_grad(j));
122     penalty = eta/h;
123     newbeta = beta(j) + penalty * this_grad;
124     beta(j) = sgn(newbeta) * fmax(0.0, fabs(newbeta) - penalty * lambda);
125 }
126 k++; // increment global counter
127 }
128 }
129

```

```
130 // (4) Update last penalties:
131 // At the very end, apply the accumulated penalty for the variables we haven't touched recently
    // . We still have to penalize some features that we did not update
132 // during the last "skip" iterations.
133 for (int j = 0; j < P; j++) {
134     double skip = k - last_update(j);
135     h = sqrt(hist_grad(j));
136     penalty = skip*eta/h;
137     beta(j) = sgn(beta(j))*fmax(0.0, fabs(beta(j)) - penalty*lambda);
138 }
139
140 return List::create(Named("alpha") = alpha,
141                     Named("beta") = beta,
142                     Named("loglik") = loglik_vec);
143 }
```

*Listing A.2: Function implementing Adagrad algorithm exploiting C++'s library Eigen.*