

SDS385 Peer evaluation 1:

Reviewee: <https://github.com/JKelle/SDS-385>

This is a summary of the suggestions I discussed with you during the meeting we had on Sept. 14.

1. Code efficiency.

- a. My suggestion is to use R's built-in function *crossprod* whenever you can. It can help preventing from doing a lot of useless computation, especially in the case of symmetric matrices. It is evident in the weighted least squares problem.
- b. Apart from that, we used the same tricks to multiply a diagonal matrix by a full matrix, which explains why our benchmarked performances look pretty similar and fast.

2. Functions structure.

- a. In my experience, it is a good programming rule to feed the functions all of the parameters of a given problem. In particular, I suggest you to feed it also with the tolerance threshold and with the maximum number of iterations instead of defining those variables inside the *gradient* function. In this way you can "play" with those parameters to see if anything changes.

3. Check performance.

- a. Apart from checking the convergence of the log-likelihood, I would also compare the final regression parameters with the ones obtained by R's built-in *glm* function. In this way you have a benchmark you can use to compare the results given by several algorithms.
- b. As far as the convergence check is concerned, you can use threshold on the relative decrement of the log-likelihood function, that is, $(lik.prec - lik.act) / lik.prec$.
- c. I like the idea of using the accuracy of the logit classifier as a performance comparison between those algorithms. They are likely to give the same result, but this is a good way to check if the logit model is appropriate.

4. Other comments.

- a. You compute the vectorized version of the log-likelihood, which is the most efficient version.
- b. You already implemented line search methods for the gradient descent, which is great!